

The TTC 2019 Live Case: BibTeX to DocBook

Antonio García-Domínguez
Aston University
B4 7ET, Birmingham, United Kingdom
a.garcia-dominguez@aston.ac.uk

Georg Hinkel
Tecan Software Competence Center
55252, Mainz-Kastel, Germany
georg.hinkel@tecan.com

Abstract

The initial transformation of a model into another model is only the first step. After the creation of the target model, it may be manually changed and the consistency with the source model may be lost or obscured. Ideally, transformation tools should have a way to check the degree of consistency between the source model and the current version of the destination model. This case presents such a scenario for a small transformation, with an automated mutation tool which will introduce changes that may or may not impact consistency. The aim of this case is to evaluate the speed and verbosity of the inter-model consistency checking in the state of the art.

1 Introduction

This live case is based on the original ATL Zoo [2] BibTeX to DocBook transformation, where a simplified version of the BibTeX reference manager's data model (shown in Figure 1) is transformed to a simplification of the DocBook document typesetting tool (shown in Figure 2).

The transformation consists of these mappings:

- From the root BIBTEXFILE, a DOCBOOK is created. The DOCBOOK has a BOOK with an ARTICLE titled “BibTeXML to DocBook”, which has in turn four SECT1 instances: “References List”, “Author List”, “Titles List” and “Journals List”.
- Any BibTeX AUTHOR is listed in the “Authors List” as a PARA.
- Any BIBTEXENTRY is listed in the “References List”, including its unique identifier and any available information.
- For each TITLEDENTRY, a PARA with its title is added to the “Titles List”
- For each BibTeX ARTICLE, a PARA with its journal name is added to the “Journals List”.
- The “Author List”, “Titles List”, and “Journals List” are all sorted lexicographically in ascending order.

The transformation is somewhat different in that it involves some sorting. The original ATL-based implementation worked well for instances with 10, 100 and 1000 entries, but seems to struggle with instances with 10000 entries.

Copyright © by the paper's authors. Use permitted under Creative Commons License Attribution 4.0 International (CC BY 4.0).

In: A. Editor, B. Coeditor (eds.): Proceedings of the XYZ Workshop, Location, Country, DD-MMM-YYYY, published at <http://ceur-ws.org>

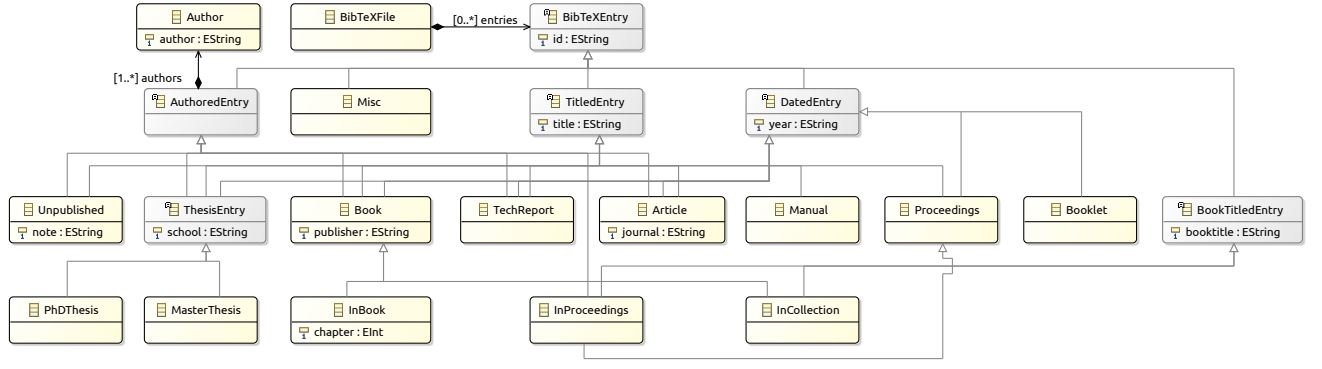


Figure 1: Class diagram for the input BibTeXXML metamodel




Figure 2: Class diagram for the target DocBook metamodel

Still, a more interesting problem is the case when the target model continues to be worked on after the initial transformation. The DocBook document would be given to an editor, which would reorganise sections, add paragraphs in the middle with further commentary and perhaps extend some of the text itself. The task is to ensure in an efficient manner that these manual editions have not impacted the consistency of the DocBook model with the original BibTeX model.

Ideally, the transformation tool or the editing environment would give some infrastructure to tackle this. However, for many tools, the approach seems to be only feasible through the creation of an external consistency checker. This case is to evaluate the current state of the art in out-of-the-box after-the-fact consistency checking. To do so, the case provides a generator which can produce source models of arbitrary size, a repackaged version of the original ATL transformation, and a mutator which can make a number of changes to the model, some of which may impact consistency. These resources were made available on Github¹ before the start of STAF 2019.

The rest of the document is structured as follows: Section 2 describes the structure of the live case. Section 3 describes the proposed tasks for this live case. Section 4 mentions the benchmark framework for those solutions that focus on raw performance. Finally, Section 5 mentions an outline of the initial audience-based evaluation across all solutions, and the approach that will be followed to derive additional prizes depending on the attributes targeted by the solutions.

2 Case Structure

The case is intended to review the different approaches for checking after-the-fact inter-model consistency between a BibTeX model and a DocBook model. The process is roughly as follows:

1. The BibTeX model is generated randomly to a certain size, by the included *generator* in the `models/generator.jar` JAR. The generator uses a Java port of the Ruby Faker² library to produce pseudorandom data given a seed. A number of random models (sizes 10, 100, 1000 and 10000) were generated in advance and included in the case resources.
2. The BibTeX model is transformed automatically to DocBook by the repackaged version of the original ATL transformation in the `bibtex2docbook.jar` JAR. This transformation deals well with models up to 1000 entries, but struggles with larger models due to recomputation of intermediate results.

¹<https://github.com/TransformationToolContest/ttc2019-live>

²<https://github.com/DiUS/java-faker>

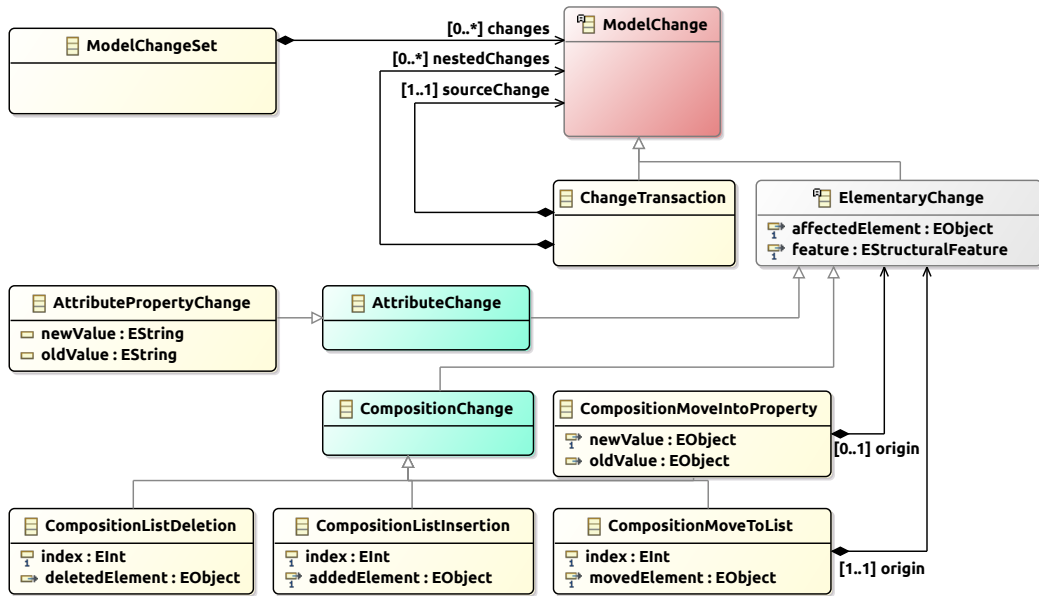


Figure 3: Class diagram for the used subset of the Changes metamodel

3. The DocBook model is edited, in this case with the automated *mutator* in the `models/mutator.jar` JAR. The mutator will operate on a DocBook file, creating a set of folders whose path will start with the specified prefix, adding -N from 1 to `nMutants`. Each folder will contain the mutated DocBook model, as well as a *change model* explaining what was done to the model: Figure 3 shows that such models have a `MODELCHANGESSET` as the root, with a number of `MODELCHANGE` instances of various types. The DocBook model will have gone through a number of random mutations according to a seed: if unspecified, the seed will be based on the current system time.

The mutator has a number of predefined *mutation operators* that will modify the model:

- Swapping paragraphs: this should break consistency in terms of sorting, but it will not result in missing information.
- Swapping sections should not break consistency.
- Deleting paragraphs/sections should break consistency.
- Appending text to a paragraph should not break consistency.
- Adding a new paragraph: unless it happens to match one of the authors, titles, or journals in its `SECT1`, it should not break consistency.

The mutated models were created in advance before the contest. There are three sets of mutated models from the generated 10/100/1000-entry models: one with a single mutation, one with two mutations, and one with three mutations.

4. A *consistency checker* would take any combination of the previous artifacts (source BibTeX, fresh DocBook, mutated DocBook, change model) and make a judgment about whether the mutated DocBook is still consistent or not.

If issues are found, it should point to the element in the source model which lacks a proper mapping on the other side, or the element in the target model which is not mapped correctly from the source model (e.g. it is not sorted anymore).

The case resources include a set of expected results from the reference EVL consistency checker. However, the concrete definition of the consistency requirements has proven to be trickier to formalize than expected. This was raised by several case authors, and in fact one of the solutions considered creating a DSL for expressing inter-model consistency.

3 Task Description

The case had an optional and a mandatory task:

- The optional task was to re-implement or improve the original transformation itself, in a way that lent itself better to after-the-fact consistency checking. A transformation tool may have better support for this, or ATL could be made to deal better with larger versions of this model.
- The mandatory task was to check for the consistency of the source BibTeX against the mutated DocBook models in the `models` directory, and report this information as efficiently and clearly as possible to the user. Ideally, this should be possible without a full re-run of the transformation. To be considered for the performance-related awards, solutions had to use the benchmarking framework in Section 4.

A reference solution based on the Epsilon Validation Language was provided. This implementation did not re-run any transformations: instead, it did a two-way consistency validation by checking from the BibTeX BIBTEXFILE, AUTHOR, BIBTEXENTRY, TITLEDENTRY, and ARTICLE types, and from the DocBook PARA types. The implementation required 98 lines of EVL code and 113 lines of Java framework integration code, and does not use the change models at all.

Solutions could focus on efficiency, conciseness, or clarity of presentation to the user. Solutions that can operate straight from the definition of the transformation (i.e. without a separate consistency checker) would be preferred. The call for solutions also invited solution authors to consider other desirable attributes, e.g. verifiability.

4 Benchmark Framework

If focusing on performance, the solution authors had to integrate their solution with the provided benchmark framework. It is based on the framework in the TTC 2017 Smart Grid case [1], and supports the automated build and execution of solutions. The benchmark consisted of three phases:

1. **Initialization**, which involved setting up the basic infrastructure (e.g. loading metamodels). These measurements are optional.
2. **Load**, which loaded the input models.
3. **Run**, which found the consistency violations in the mutated DocBook model.

4.1 Solution requirements

Each solution had to print to the standard output a line with the following fields, separated by semicolons (“;”):

- **Tool**: name of the tool.
- **MutantSet**: set of mutants used (“single”, “double” or “triple”).
- **Source**: base name of the input BibTeX model (e.g. “random10.bibtex”).
- **Mutant**: integer starting at 1, identifying the mutant model within this set.
- **RunIndex**: index of the run of this combination of tools and inputs.
- **PhaseName**: name of the phase being run.
- **MetricName**: the name of the metric. It may be the used **Memory** in bytes, the wall clock **Time** spent in integer nanoseconds, or the number of consistency **Problems** found in the mutated DocBook model.

To enable automatic execution by the benchmark framework, solutions were in the form of a subfolder within the `solutions` folder of the main repository³, with a `solution.ini` file stating how the solution should be built and how it should be run. As an example, the `solution.ini` file for the reference solution is shown on Listing 1. In the `build` section, the `default` option specifies the command to build and test the solution, and

³<https://github.com/TransformationToolContest/ttc2019-live>

Listing 1: `solution.ini` file for the reference Epsilon solution

```

1 [build]
2 default=true
3 skipTests=true
4
5 [run]
6 cmd=JAVA_OPTS="-Xms4g" java -jar epsilon.jar

```

the `skipTests` option specifies the command to build the solution while skipping unit tests. In the `run` section, the `cmd` option specifies the command to run the solution.

The repetition of executions as defined in the benchmark configuration was done by the benchmark. For 5 runs, the specified command will be called 5 times, passing any required information (e.g. run index, or input model name) through environment variables. Solutions could not save intermediate data between different runs: each run had to be entirely independent.

The name and absolute path of the input model, the run index and the name of the tool were passed using environment variables `Tool`, `MutantSet`, `SourcePath`, `Mutant`, `MutantPath`, and `RunIndex`.

4.2 Running the benchmark

The benchmark framework only required Python 3.3 to be installed. Solutions could use any languages or frameworks, as long as they could run without human input. Since all the performance-oriented solutions this year were compatible with GNU/Linux, it was possible for the case authors to create a `Dockerfile` with all solutions built in, for the sake of reproducibility. The resulting image is available on Docker Hub⁴, and it is automatically rebuilt on any push to the repository.

If all prerequisites are fulfilled, the benchmark can be run using Python with the command `python scripts/run.py`. Additional options can be queried using the option `--help`. The benchmark framework can be configured through the `config/config.json` file: this includes the input models to be evaluated (some of which have been excluded by default due to their high cost with the sample solution), the names of the tools to be run, the number of runs per tool+model, and the timeout for each command in milliseconds.

5 Evaluation

The evaluation operated on several dimensions:

- How efficient was the approach in time and space (memory)? The reference ATL solution struggled with large models, and the reference solution was not been designed with performance in mind.
- Was consistency checking directly supported by the transformation approach? Many tools lack this capability, though it might be interesting as an additional execution mode if the target model has seen manual changes since it was generated.
- How informative and accessible was the feedback that can be provided by the approach?

Authors were invited to make submissions targeting other quality attributes.

References

- [1] Georg Hinkel. The TTC 2017 Outage System Case for Incremental Model Views. In *Proceedings of the 10th Transformation Tool Contest*, volume 2026, pages 3–12, Marburg, Germany, July 2017. CEUR-WS.org.
- [2] Guillaume Savaton. BibTeXML to DocBook, ATL Transformations. <https://www.eclipse.org/atl/atlTransformations/#BibTeXML2DocBook>, February 2006. Last accessed on 2019-07-15. Archived on <http://archive.is/HdoHM>.

⁴<https://hub.docker.com/r/bluezio/ttc2019-live-git>